

lab_1

October 31, 2023

Aim: Predict the price of the Uber ride from a given pickup point to the agreed drop-off location.

Name: Yash Ghorpade

Div: BE-A

Roll No: B211046

```
[2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
```

```
[3]: df = pd.read_csv('uber - uber.csv')
df.info()
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 200000 entries, 0 to 199999

Data columns (total 9 columns):

#	Column	Non-Null Count	Dtype
0	Unnamed: 0	200000 non-null	int64
1	key	200000 non-null	object
2	fare_amount	200000 non-null	float64
3	pickup_datetime	200000 non-null	object
4	pickup_longitude	200000 non-null	float64
5	pickup_latitude	200000 non-null	float64
6	dropoff_longitude	199999 non-null	float64
7	dropoff_latitude	199999 non-null	float64
8	passenger_count	200000 non-null	int64

dtypes: float64(5), int64(2), object(2)

memory usage: 13.7+ MB

```
[4]: df.shape
```

```
[4]: (200000, 9)
```

```
[5]: df.head()
```

```
[5]: Unnamed: 0      key  fare_amount  pickup_datetime \
0    24238194  2015-05-07 19:52:06          7.5  2015-05-07 19:52:06 UTC
```

1	27835199	2009-07-17 20:04:56	7.7	2009-07-17 20:04:56 UTC
2	44984355	2009-08-24 21:45:00	12.9	2009-08-24 21:45:00 UTC
3	25894730	2009-06-26 8:22:21	5.3	2009-06-26 08:22:21 UTC
4	17610152	2014-08-28 17:47:00	16.0	2014-08-28 17:47:00 UTC

	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude\ 0
	-73.999817	40.738354	-73.999512	40.723217
1	-73.994355	40.728225	-73.994710	40.750325
2	-74.005043	40.740770	-73.962565	40.772647
3	-73.976124	40.790844	-73.965316	40.803349
4	-73.925023	40.744085	-73.973082	40.761247

	passenger_count
0	1
1	1
2	1
3	3
4	5

```
[6]: df.head(7)
```

```
[6]: Unnamed: 0      key  fare_amount      pickup_datetime \
0  24238194  2015-05-07 19:52:06         7.5  2015-05-07 19:52:06 UTC
1  27835199  2009-07-17 20:04:56         7.7  2009-07-17 20:04:56 UTC
2  44984355  2009-08-24 21:45:00        12.9  2009-08-24 21:45:00 UTC
3  25894730  2009-06-26 8:22:21         5.3  2009-06-26 08:22:21 UTC
4  17610152  2014-08-28 17:47:00        16.0  2014-08-28 17:47:00 UTC
5  44470845  2011-02-12 2:27:09         4.9  2011-02-12 02:27:09 UTC
6  48725865  2014-10-12 7:04:00        24.5  2014-10-12 07:04:00 UTC
```

	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude \
0	-73.999817	40.738354	-73.999512	40.723217
1	-73.994355	40.728225	-73.994710	40.750325
2	-74.005043	40.740770	-73.962565	40.772647
3	-73.976124	40.790844	-73.965316	40.803349
4	-73.925023	40.744085	-73.973082	40.761247
5	-73.969019	40.755910	-73.969019	40.755910
6	-73.961447	40.693965	-73.871195	40.774297

	passenger_count
0	1
1	1
2	1
3	3
4	5
5	1
6	5

```
[7]: df.isnull()
```

```
[7]:      Unnamed: 0      key fare_amount pickup_datetime pickup_longitude \
0          False False          False          False          False
1          False False          False          False          False
2          False False          False          False          False
3          False False          False          False          False
4          False False          False          False          False
...
199995      False False          False          False          False
199996      False False          False          False          False
199997      False False          False          False          False
199998      False False          False          False          False
199999      False False          False          False          False

      pickup_latitude dropoff_longitude dropoff_latitude passenger_count
0          False          False          False          False
1          False          False          False          False
2          False          False          False          False
3          False          False          False          False
4          False          False          False          False
...
199995      False          False          False          False
199996      False          False          False          False
199997      False          False          False          False
199998      False          False          False          False
199999      False          False          False          False
```

[200000 rows x 9 columns]

```
[8]: df.dropna(inplace=True)
df.drop(columns=["Unnamed: 0", "key"], inplace=True)
df.head()
```

```
[8] :      fare_amount      pickup_datetime pickup_longitude pickup_latitude \
0          7.5  2015-05-07 19:52:06 UTC      -73.999817      40.738354
1          7.7  2009-07-17 20:04:56 UTC      -73.994355      40.728225
2         12.9  2009-08-24 21:45:00 UTC      -74.005043      40.740770
3          5.3  2009-06-26 08:22:21 UTC      -73.976124      40.790844
4         16.0  2014-08-28 17:47:00 UTC      -73.925023      40.744085

      dropoff_longitude dropoff_latitude passenger_count
0      -73.999512      40.723217          1
1      -73.994710      40.750325          1
2      -73.962565      40.772647          1
3      -73.965316      40.803349          3
4      -73.973082      40.761247          5
```

```
[9]: df.isnull().sum()
```

```
[9] : fare_amount      0
      pickup_datetime  0
      pickup_longitude 0
      pickup_latitude  0
      dropoff_longitude 0
      dropoff_latitude 0
      passenger_count   0
      dtype: int64
```

```
[10]: df.dtypes
```

```
[10] : fare_amount      float64
      pickup_datetime   object
      pickup_longitude  float64
      pickup_latitude   float64
      dropoff_longitude float64
      dropoff_latitude  float64
      passenger_count   int64
      dtype: object
```

```
[11]: df.pickup_datetime = pd.to_datetime(df.pickup_datetime)
      df.dtypes
```

```
[11] : fare_amount      float64
      pickup_datetime  datetime64[ns, UTC]
      pickup_longitude float64
      pickup_latitude  float64
      dropoff_longitude float64
      dropoff_latitude  float64
      passenger_count   int64
      dtype: object
```

```
[12]: df
      df = df.assign(hour = df.pickup_datetime.dt.hour,
      day = df.pickup_datetime.dt.day,
      month = df.pickup_datetime.dt.month,
      year = df.pickup_datetime.dt.year,
      dayofweek = df.pickup_datetime.dt.dayofweek)
```

```
[13]: df
```

```
[13] :      fare_amount      pickup_datetime  pickup_longitude \
0          7.5 2015-05-07 19:52:06+00:00      -73.999817
1          7.7 2009-07-17 20:04:56+00:00      -73.994355
2         12.9 2009-08-24 21:45:00+00:00      -74.005043
3          5.3 2009-06-26 08:22:21+00:00      -73.976124
4         16.0 2014-08-28 17:47:00+00:00      -73.925023
...          ...          ...          ...
```

199995	3.0	2012-10-28 10:49:00+00:00	-73.987042
199996	7.5	2014-03-14 01:09:00+00:00	-73.984722
199997	30.9	2009-06-29 00:42:00+00:00	-73.986017
199998	14.5	2015-05-20 14:56:25+00:00	-73.997124
199999	14.1	2010-05-15 04:08:00+00:00	-73.984395

	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count	\
0	40.738354	-73.999512	40.723217	1	
1	40.728225	-73.994710	40.750325	1	
2	40.740770	-73.962565	40.772647	1	
3	40.790844	-73.965316	40.803349	3	
4	40.744085	-73.973082	40.761247	5	
...	
199995	40.739367	-73.986525	40.740297	1	
199996	40.736837	-74.006672	40.739620	1	
199997	40.756487	-73.858957	40.692588	2	
199998	40.725452	-73.983215	40.695416	1	
199999	40.720077	-73.985508	40.768793	1	

	hour	day	month	year	dayofweek
0	19	7	5	2015	3
1	20	17	7	2009	4
2	21	24	8	2009	0
3	8	26	6	2009	4
4	17	28	8	2014	3
...
199995	10	28	10	2012	6
199996	1	14	3	2014	4
199997	0	29	6	2009	0
199998	14	20	5	2015	2
199999	4	15	5	2010	5

[199999 rows x 12 columns]

```
[14]: df = df.drop(["pickup_datetime"], axis =1)
df
```

```
[14]:
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	\
0	7.5	-73.999817	40.738354	-73.999512	
1	7.7	-73.994355	40.728225	-73.994710	
2	12.9	-74.005043	40.740770	-73.962565	
3	5.3	-73.976124	40.790844	-73.965316	
4	16.0	-73.925023	40.744085	-73.973082	
...	
199995	3.0	-73.987042	40.739367	-73.986525	
199996	7.5	-73.984722	40.736837	-74.006672	
199997	30.9	-73.986017	40.756487	-73.858957	

199998	14.5	-73.997124	40.725452	-73.983215
199999	14.1	-73.984395	40.720077	-73.985508

	dropoff_latitude	passenger_count	hour	day	month	year	dayofweek
0	40.723217	1	19	7	5	2015	3
1	40.750325	1	20	17	7	2009	4
2	40.772647	1	21	24	8	2009	0
3	40.803349	3	8	26	6	2009	4
4	40.761247	5	17	28	8	2014	3
...
199995	40.740297	1	10	28	10	2012	6
199996	40.739620	1	1	14	3	2014	4
199997	40.692588	2	0	29	6	2009	0
199998	40.695416	1	14	20	5	2015	2
199999	40.768793	1	4	15	5	2010	5

[199999 rows x 11 columns]

```
[16]: from math import *

def distance_formula(longitude1, latitude1, longitude2, latitude2):
    travel_dist = []

    for pos in range(len(longitude1)):


        lon1, lan1, lon2, lan2 = map(radians, [longitude1[pos], latitude1[pos],
        ↪ longitude2[pos], latitude2[pos]])
        dist_lon = lon2 - lon1
        dist_lan = lan2 - lan1
        a = sin(dist_lan/2)**2 + cos(lan1) * cos(lan2) * sin(dist_lon/2)**2

        #radius of earth = 6371
        c = 2 * asin(sqrt(a)) * 6371
        travel_dist.append(c)

    return travel_dist
```

```
[17]: df["dist_travel_km"] = distance_formula(df.pickup_longitude.to_numpy(), df.
    ↪ pickup_latitude.to_numpy(), df.dropoff_longitude.to_numpy(),df.
    ↪ dropoff_latitude.to_numpy())
```

```
[18]: df.plot(kind = "box",subplots = True,layout = (6,2),figsize=(15,20))
    #Boxplot to check the outliers
    plt.show()
```



```
[21]: def remove_outlier(df1 , col):
      Q1 = df1[col].quantile(0.25)
      Q3 = df1[col].quantile(0.75)
      IQR = Q3 - Q1
```

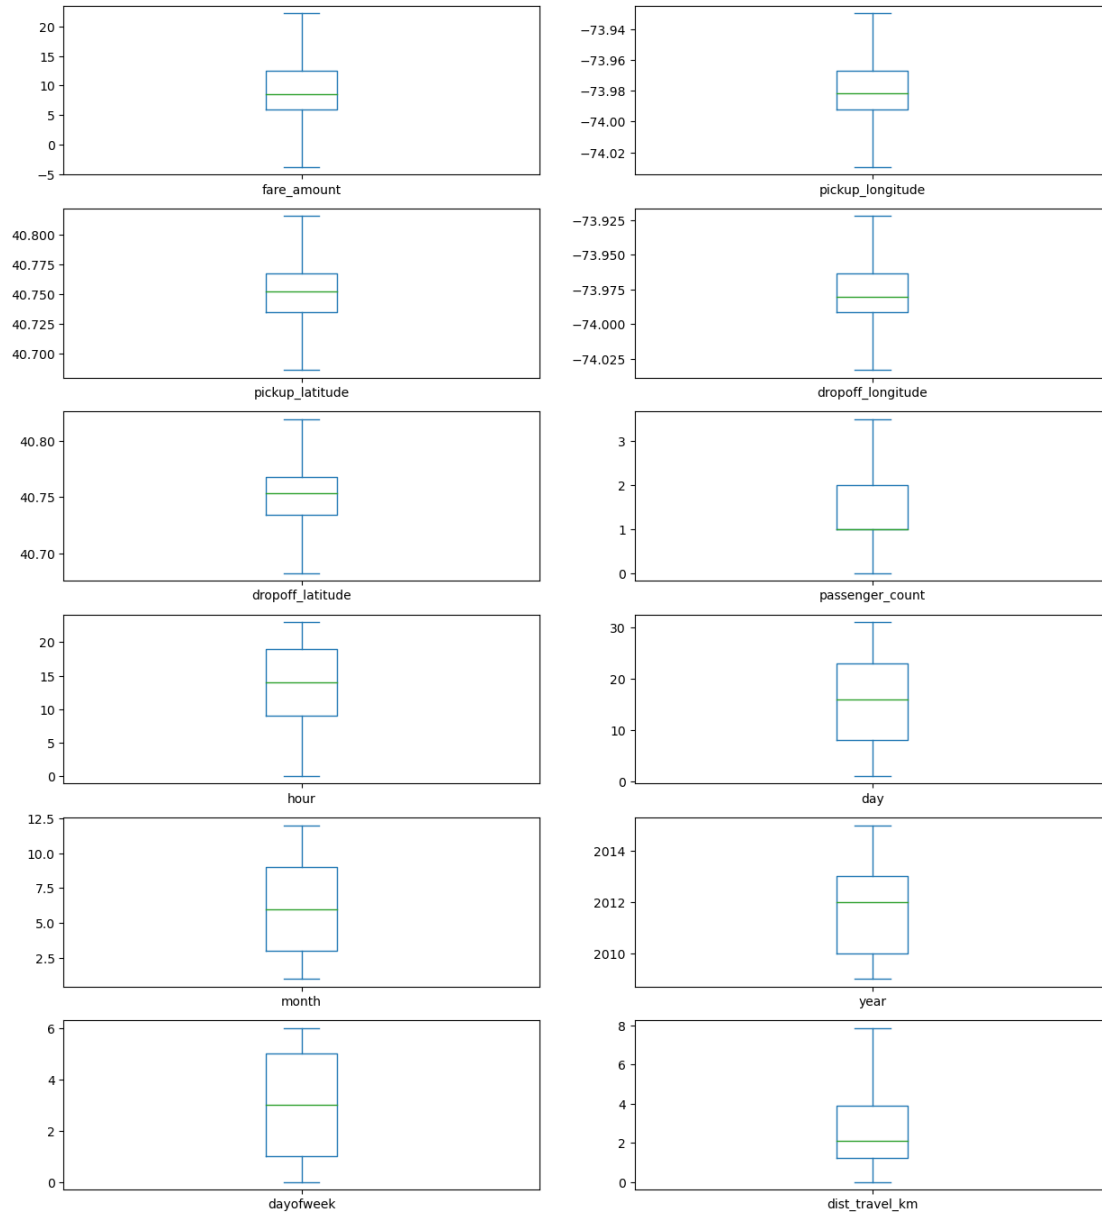


```
lower_whisker = Q1-1.5*IQR
upper_whisker = Q3+1.5*IQR
df[col] = np.clip(df1[col] , lower_whisker , upper_whisker)
return df1
```

```
def treat_outliers_all(df1 , col_list):
    for c in col_list:
        df1 = remove_outlier(df , c)
    return df1
```

```
[22]: df = treat_outliers_all(df , df.iloc[:, 0::])
```

```
[23]: df.plot(kind = "box",subplots = True,layout = (7,2),figsize=(15,20))
plt.show()
```



```
[24]: corr = df.corr()
corr
```

```
[24]:
```

	fare_amount	pickup_longitude	pickup_latitude	\
fare_amount	1.000000	0.154053	-0.110857	
pickup_longitude	0.154053	1.000000	0.259496	
pickup_latitude	-0.110857	0.259496	1.000000	
dropoff_longitude	0.218681	0.425622	0.048887	
dropoff_latitude	-0.125874	0.073313	0.515736	
passenger_count	0.015798	-0.013202	-0.012879	

hour	-0.023605	0.011590	0.029691
day	0.004552	-0.003194	-0.001544
month	0.030815	0.001168	0.001561
year	0.141271	0.010193	-0.014247
dayofweek	0.013664	-0.024645	-0.042304
dist_travel_km	0.844369	0.098074	-0.046825

	dropoff_longitude	dropoff_latitude	passenger_count \
fare_amount	0.218681	-0.125874	0.015798
pickup_longitude	0.425622	0.073313	-0.013202
pickup_latitude	0.048887	0.515736	-0.012879
dropoff_longitude	1.000000	0.245670	-0.009304
dropoff_latitude	0.245670	1.000000	-0.006329
passenger_count	-0.009304	-0.006329	1.000000
hour	-0.046560	0.019765	0.020260
day	-0.004008	-0.003498	0.002699
month	0.002392	-0.001191	0.010353
year	0.011347	-0.009595	-0.009743
dayofweek	-0.003337	-0.031932	0.048542
dist_travel_km	0.186533	-0.038873	0.009729


	hour	day	month	year	dayofweek \
fare_amount	-0.023605	0.004552	0.030815	0.141271	0.013664
pickup_longitude	0.011590	-0.003194	0.001168	0.010193	-0.024645
pickup_latitude	0.029691	-0.001544	0.001561	-0.014247	-0.042304
dropoff_longitude	-0.046560	-0.004008	0.002392	0.011347	-0.003337
dropoff_latitude	0.019765	-0.003498	-0.001191	-0.009595	-0.031932
passenger_count	0.020260	0.002699	0.010353	-0.009743	0.048542
hour	1.000000	0.004664	-0.003924	0.002162	-0.086956
day	0.004664	1.000000	-0.017358	-0.012165	0.005609
month	-0.003924	-0.017358	1.000000	-0.115860	-0.008785
year	0.002162	-0.012165	-0.115860	1.000000	0.006116
dayofweek	-0.086956	0.005609	-0.008785	0.006116	1.000000
dist_travel_km	-0.038348	0.003080	0.011626	0.024270	0.027066

	dist_travel_km
fare_amount	0.844369
pickup_longitude	0.098074
pickup_latitude	-0.046825
dropoff_longitude	0.186533
dropoff_latitude	-0.038873
passenger_count	0.009729
hour	-0.038348
day	0.003080
month	0.011626
year	0.024270
dayofweek	0.027066

dist_travel_km 1.000000

```
[25]: fig,axis = plt.subplots(figsize = (10,6))
sns.heatmap(df.corr(),annot = True)
```

[25]: <Axes: >



```
[26]: df_x = _
      ↪ df[["pickup_longitude","pickup_latitude","dropoff_longitude","dropoff_latitude","passenger_
df_y = df["fare_amount"]
```

```
[27]: x_train, x_test, y_train, y_test = train_test_split(df_x, df_y, test_size=0.46, _
      ↪ random_state=46)
```

```
[28]: df
```

```
[28]:   fare_amount  pickup_longitude  pickup_latitude  dropoff_longitude \
0          7.50      -73.999817      40.738354      -73.999512
1          7.70      -73.994355      40.728225      -73.994710
2         12.90      -74.005043      40.740770      -73.962565
3          5.30      -73.976124      40.790844      -73.965316
4         16.00      -73.929787      40.744085      -73.973082
```

```

...
199995      3.00      -73.987042      40.739367      -73.986525
199996      7.50      -73.984722      40.736837      -74.006672
199997     22.25      -73.986017      40.756487      -73.922036
199998     14.50      -73.997124      40.725452      -73.983215
199999     14.10      -73.984395      40.720077      -73.985508

      dropoff_latitude  passenger_count  hour  day  month  year  dayofweek  \
0          40.723217              1.0   19    7     5  2015           3
1          40.750325              1.0   20   17     7  2009           4
2          40.772647              1.0   21   24     8  2009           0
3          40.803349              3.0    8   26     6  2009           4
4          40.761247              3.5   17   28     8  2014           3

...
199995      40.740297              1.0   10   28    10  2012           6
199996      40.739620              1.0    1   14     3  2014           4
199997      40.692588              2.0    0   29     6  2009           0
199998      40.695416              1.0   14   20     5  2015           2
199999      40.768793              1.0    4   15     5  2010           5

      dist_travel_km
0          1.683323
1          2.457590
2          5.036377
3          1.661683
4          4.475450

...
199995      0.112210
199996      1.875050
199997      7.865091
199998      3.539715
199999      5.417783

```

[199999 rows x 12 columns]

```
[29]: from sklearn.linear_model import LinearRegression
```

```
[30]: reg = LinearRegression()
```

```
[31]: reg.fit(x_train, y_train)
```

```
[31]: LinearRegression()
```

```
[32]: y_pred_lin = reg.predict(x_test)
      print(y_pred_lin)
```

```
[ 9.71080842  7.93130063  9.03773769 ... 10.20762805  9.54503729
  5.97704907]
```

```
[33]: from sklearn.ensemble import RandomForestRegressor
```

```
[38]: rf = RandomForestRegressor(n_estimators=10)
      rf.fit(x_train,y_train)
```

```
[38]: RandomForestRegressor(n_estimators=10)
```

```
[39]: y_pred_rf = rf.predict(x_test)
      print(y_pred_rf)
```

```
[11.65  5.65  8.85 ...  8.9  10.48  5.35]
```

```
[40]: cols = ["Model", "RMSE", "R-Squared"]
```

```
[41]: result_tabulation = pd.DataFrame(columns = cols)
```

```
[42]: from sklearn import metrics
      from sklearn.metrics import r2_score
      reg_RMSE = np.sqrt(metrics.mean_squared_error(y_test, y_pred_lin))
      reg_squared = r2_score(y_test, y_pred_lin)
      full_metrics = pd.Series({"Model": "Linear Regression", "RMSE" : reg_RMSE,
      ↪ "R-Squared" : reg_squared})
```

```
[43]: result_tabulation = result_tabulation._append(full_metrics, ignore_index = True)
```

```
[44]: result_tabulation
```

```
[44]:
```

	Model	RMSE	R-Squared
0	Linear Regression	2.741174	0.744953

```
[45]: rf_RMSE = np.sqrt(metrics.mean_squared_error(y_test, y_pred_rf))
      rf_squared = r2_score(y_test, y_pred_rf)
      full_metrics = pd.Series({"Model": "Random Forest ", "RMSE":rf_RMSE,
      ↪ "R-Squared": rf_squared})
```

```
[46]: result_tabulation = result_tabulation._append(full_metrics, ignore_index = True)
```

```
[47]: result_tabulation
```

```
[47]:
```

	Model	RMSE	R-Squared
0	Linear Regression	2.741174	0.744953
1	Random Forest	2.526872	0.783272

```
[ ]:
```